

Bayesian Parameter Extraction

Benjamin Audren

Institute of Theoretical Physics
École Polytechnique Fédérale de Lausanne

27/10/2014

Outline

1 Introduction

- Need for a statistical tool
- Bayesian approach

2 Metropolis Hastings

- Foreword
- Algorithm
- Importance Sampling

Outline

1 Introduction

- Need for a statistical tool
- Bayesian approach

2 Metropolis Hastings

Bibliography

Book

- Bayesian methods in cosmology (Hobson, Jaffe, Liddle, Mukherjee and Parkinson)

Review Articles

- Comparison of sampling techniques for Bayesian parameter estimation (*Rupert Allison, Joanna Dunkley, arXiv:1308.2675*)
- Bayes in the sky: Bayesian inference and model selection in cosmology (*Roberto Trotta, arXiv:0803.4089*)

Precision Cosmology

Precise measurement of random realisations

- Beginning of **precision** experiments with **WMAP**

Precision Cosmology

Precise measurement of random realisations

- Beginning of **precision** experiments with **WMAP**
- Cosmology is a **stochastic theory**, requiring observation of **statistical quantities**

Precision Cosmology

Precise measurement of random realisations

- Beginning of **precision** experiments with **WMAP**
- Cosmology is a **stochastic theory**, requiring observation of **statistical quantities**
- A **set of parameters** produce prediction for these **statistical quantities**

Precision Cosmology

Precise measurement of random realisations

- Beginning of **precision** experiments with **WMAP**
- Cosmology is a **stochastic theory**, requiring observation of **statistical quantities**
- A **set of parameters** produce prediction for these **statistical quantities**
- How to **infer the value of the parameters from the data ?**

Precision Cosmology

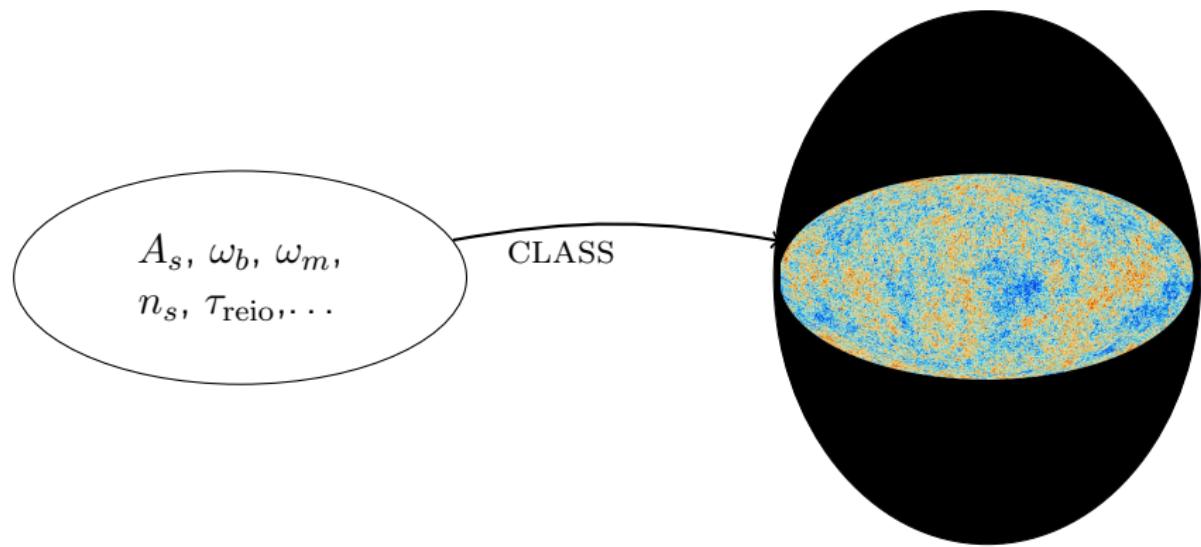
Precise measurement of random realisations

- Beginning of **precision** experiments with **WMAP**
- Cosmology is a **stochastic theory**, requiring observation of **statistical quantities**
- A **set of parameters** produce prediction for these **statistical quantities**
- How to **infer the value of the parameters from the data ?**

Similar to Particle Physics

Prediction only for, e.g. the **rate of decay** of a particle. Information acquired when statistically observing this decay channel (**how many times did it decay to this particular product ?**)

The big picture



Bayesian approach

from Bayesian Methods in Cosmology

Bayesian methods

- All quantities are considered **statistical**

Bayesian approach

from Bayesian Methods in Cosmology

Bayesian methods

- All quantities are considered **statistical**
- The question we can answer is: **is this model better than another ?**

Bayesian approach

from Bayesian Methods in Cosmology

Bayesian methods

- All quantities are considered **statistical**
- The question we can answer is: **is this model better than another ?**
- We infer **credible regions** in which, given a model, the parameters live.

Bayesian approach

from Bayesian Methods in Cosmology

Bayesian methods

- All quantities are considered **statistical**
- The question we can answer is: **is this model better than another ?**
- We infer **credible regions** in which, given a model, the parameters live.
- Our knowledge depends on the **data measured**.

Bayesian approach

Definitions

Given a **context** \mathcal{I} and a set of **data** D :

- θ : **continuous** values for a parameter set
 - $\text{pr}(\theta) \geq 0$: **probability** function is positive
 - $\int \text{pr}(\theta) d\theta = 1$: **Sum rule**
 - $\text{pr}(\phi, \theta) = \text{pr}(\phi|\theta)\text{pr}(\theta)$: **Product rule**
- $\left. \right\} \parallel \mathcal{I}$

Bayes Theorem

Bayes Theorem

All quantities are given in the context \mathcal{I}

$$\begin{array}{rclcrcl} \text{pr}(\theta)\text{pr}(D|\theta) & = & \text{pr}(\theta, D) & = & \text{pr}(D)\text{pr}(\theta|D) \\ \textbf{Prior} \times \textbf{Likelihood} & = & \textbf{Joint} & = & \textbf{Evidence} \times \textbf{Posterior} \\ \pi(\theta)\mathcal{L}(\theta) & = & \dots & = & E\mathcal{P}(\theta) \\ \text{Input} & & \longrightarrow & & \text{Output} \end{array}$$

Bayes Theorem

Bayes Theorem

All quantities are given in the context \mathcal{I}

$$\begin{array}{lclclcl} \text{pr}(\theta) \text{pr}(D|\theta) & = & \text{pr}(\theta, D) & = & \text{pr}(D) \text{pr}(\theta|D) \\ \textbf{Prior} \times \text{Likelihood} & = & \text{Joint} & = & \text{Evidence} \times \text{Posterior} \\ \pi(\theta) \mathcal{L}(\theta) & = & \dots & = & E\mathcal{P}(\theta) \\ \text{Input} & & \longrightarrow & & \text{Output} \end{array}$$

Prior

A priori information on the parameters. Most of the time, it is **flat** (uniform chance to be inside a given volume).

Bayes Theorem

Bayes Theorem

All quantities are given in the context \mathcal{I}

$$\begin{array}{lclclcl} \text{pr}(\theta) \text{pr}(\mathbf{D}|\theta) & = & \text{pr}(\theta, D) & = & \text{pr}(D) \text{pr}(\theta|D) \\ \text{Prior} \times \text{Likelihood} & = & \text{Joint} & = & \text{Evidence} \times \text{Posterior} \\ \pi(\theta) \mathcal{L}(\theta) & = & \dots & = & E\mathcal{P}(\theta) \\ \text{Input} & & \longrightarrow & & \text{Output} \end{array}$$

Likelihood

Given by the instrument operated at *known input*. If uncontrolled unknown: **nuisance parameters**.

Bayes Theorem

Bayes Theorem

All quantities are given in the context \mathcal{I}

$$\begin{array}{ccccccccc} \text{pr}(\theta)\text{pr}(D|\theta) & = & \text{pr}(\theta, D) & = & \text{pr}(\mathbf{D})\text{pr}(\theta|D) \\ \text{Prior} \times \text{Likelihood} & = & \text{Joint} & = & \text{Evidence} \times \text{Posterior} \\ \pi(\theta)\mathcal{L}(\theta) & = & \dots & = & \mathbf{EP}(\theta) \\ \text{Input} & & \longrightarrow & & \text{Output} \end{array}$$

Evidence

Only recovered with Nested Sampling, gives an information on how well the given context suits the data.

Bayes Theorem

Bayes Theorem

All quantities are given in the context \mathcal{I}

$$\begin{array}{lclclcl} \text{pr}(\theta)\text{pr}(D|\theta) & = & \text{pr}(\theta, D) & = & \text{pr}(D)\text{pr}(\theta|\mathbf{D}) \\ \text{Prior} \times \text{Likelihood} & = & \text{Joint} & = & \text{Evidence} \times \text{Posterior} \\ \pi(\theta)\mathcal{L}(\theta) & = & \dots & = & E\mathcal{P}(\theta) \\ \text{Input} & & \longrightarrow & & \text{Output} \end{array}$$

Posterior Distribution

The name of the game. Inferred distribution of probability, after using the data.

Bayes Theorem

Complete Rules

$$\int \pi(\theta) d\theta = \int \mathcal{P}(\theta) d\theta = 1$$

$$E = \int \mathcal{L}(\theta) \pi(\theta) d\theta$$

$$\mathcal{P}(\theta) = \frac{\pi(\theta) \mathcal{L}(\theta)}{E}$$

Bayesian approach: issues

What are the problems ?

- **Evidence** is hard to compute because...
- we don't usually know the **Likelihood** function (analytically).
- so we don't know where to **sample** it (many dimensions)...
- and it might be **computationally expensive**.

Bayesian approach: issues

What are the problems ?

- **Evidence** is hard to compute because...
- we don't usually know the **Likelihood** function (analytically).
- so we don't know where to **sample** it (many dimensions)...
- and it might be **computationally expensive**.

How to deal with it ?

- we have to sample **randomly** the volume (*all methods*)
- we can avoid computing this integral by **computing ratios** (*mcmc*)

Methods that give the Evidence

Comparison between models

Evidence is **how well your context explains the data**. Nested Sampling gives this. Others dont... but best-fit likelihood gives some indication.

Methods that give the Evidence

Comparison between models

Evidence is **how well your context explains the data**. Nested Sampling gives this. Others don't... but best-fit likelihood gives some indication.

Caution

These methods will only give you a **probable** answer, not a definite one. It is the price you pay for not doing your integral exactly.

But...

Living without the evidence

Theoretically motivated model: we want to know the values of this parameter to explain the data. Then, the **posterior** is an interesting quantity, and the **evidence** can be left aside temporarily.

But...

Living without the evidence

Theoretically motivated model: we want to know the values of this parameter to explain the data. Then, the **posterior** is an interesting quantity, and the **evidence** can be left aside temporarily.

Nonetheless

Beware of the best-fit likelihood value !

Outline

1 Introduction

2 Metropolis Hastings

- Foreword
- Algorithm
- Importance Sampling

Algorithm and Sampler

Foreword

Algorithm

Tells you how to **choose** which point you move, and if you **accept** a new point or not

Sampler

Tells you how to **move**, **select a** new point.

Sometimes used **interchangeably**

Algorithm and Sampler

Foreword

問手杯尊

available in Monte Python

As of v2.0.0, you can use **MultiNest** (nested sampling, by Farhan Feroz & Mike Hobson), and the **CosmoHammer** (emcee, by Joel Akeret & Sebastian Seehars)

Algorithm and Sampler

Foreword



available in Monte Python

As of v2.0.0, you can use **MultiNest** (nested sampling, by Farhan Feroz & Mike Hobson), and the **CosmoHammer** (emcee, by Joel Akeret & Sebastian Seehars)

MultiNest

MultiNest was easily implemented in Monte Python thanks to **PyMultiNest**, done by Johannes Buchner, and the help of Jesus Torrado

Algorithm and Sampler

Foreword



available in Monte Python

As of v2.0.0, you can use **MultiNest** (nested sampling, by Farhan Feroz & Mike Hobson), and the **CosmoHammer** (emcee, by Joel Akeret & Sebastian Seehars)

MultiNest

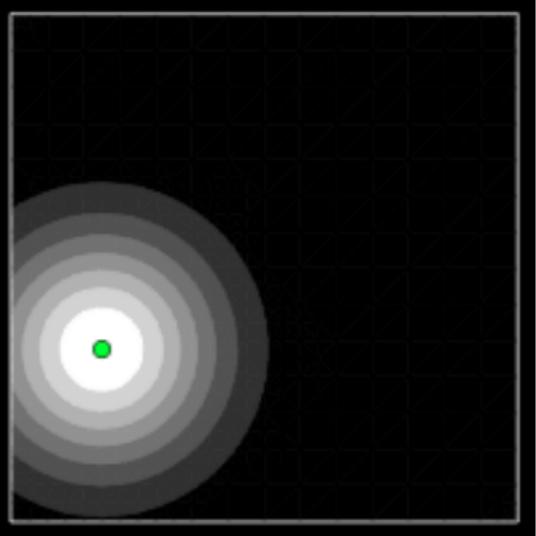
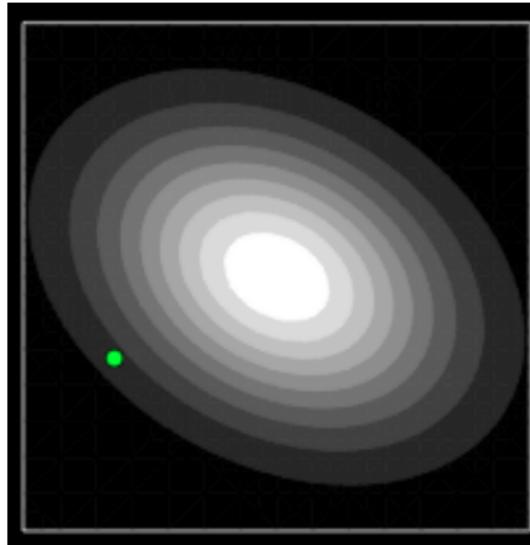
MultiNest was easily implemented in Monte Python thanks to **PyMultiNest**, done by Johannes Buchner, and the help of Jesus Torrado

Cosmo Hammer

Thanks to Joel and Sebastian for helping setting this up

Metropolis-Hastings

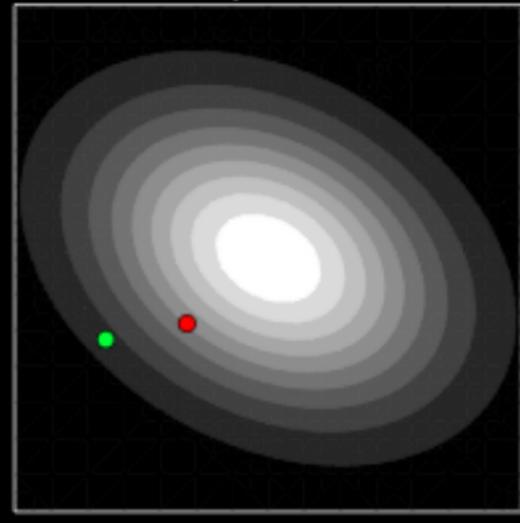
Courtesy of Sebastian Seehars



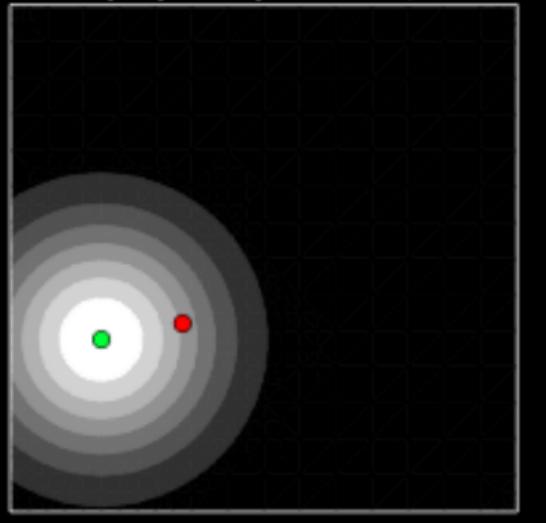
Metropolis-Hastings

Courtesy of Sebastian Seehars

initial position θ_0

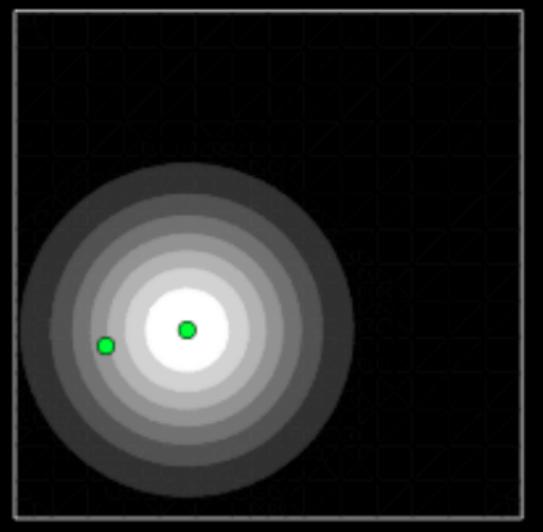
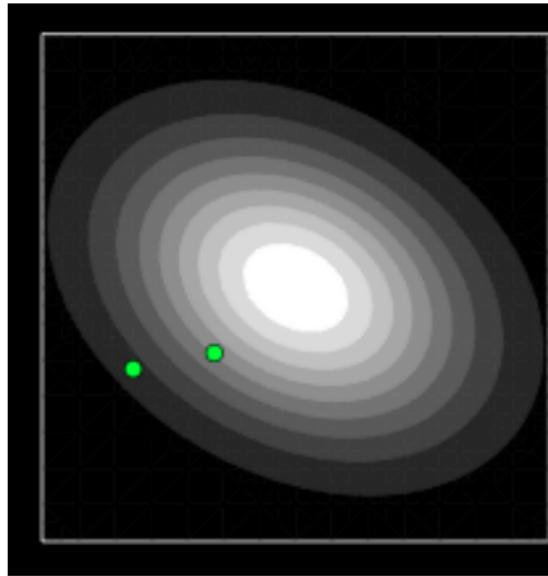


proposed position θ_p



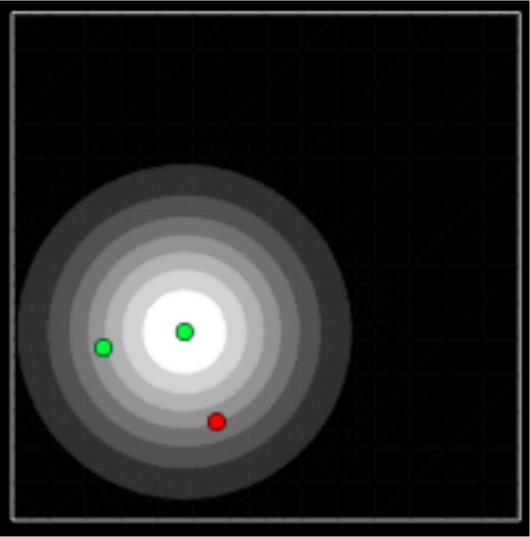
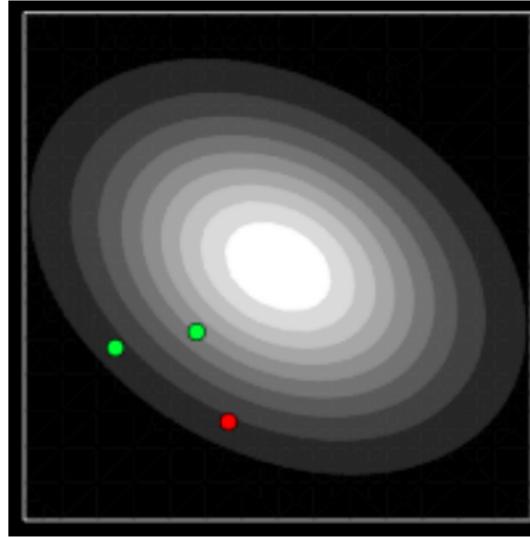
Metropolis-Hastings

Courtesy of Sebastian Seehars



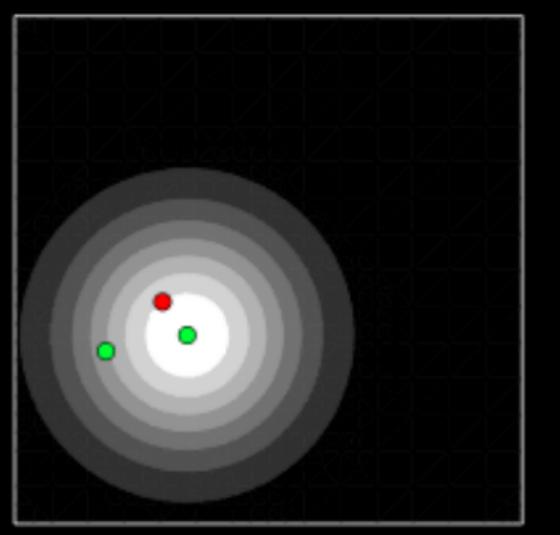
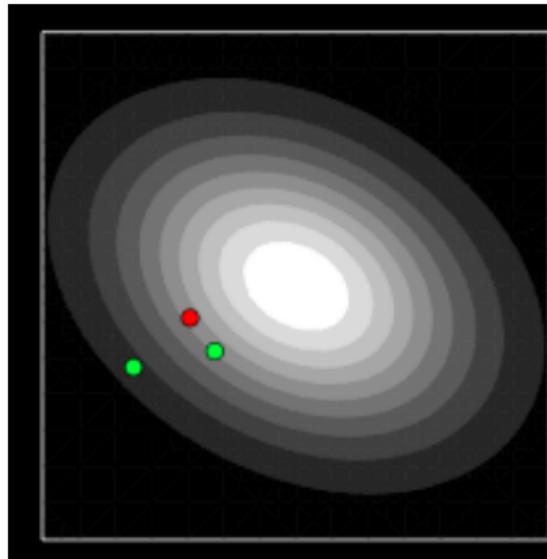
Metropolis-Hastings

Courtesy of Sebastian Seehars



Metropolis-Hastings

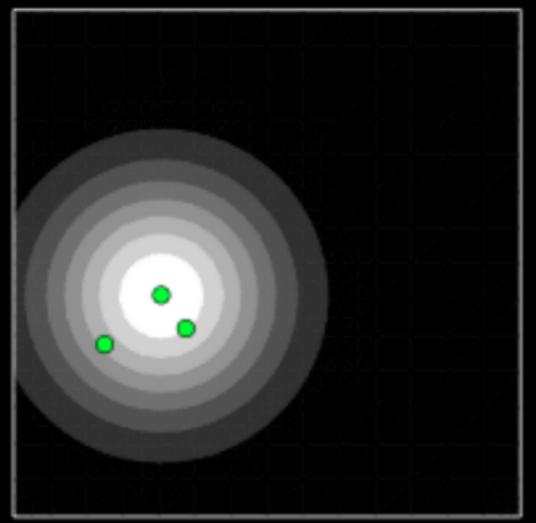
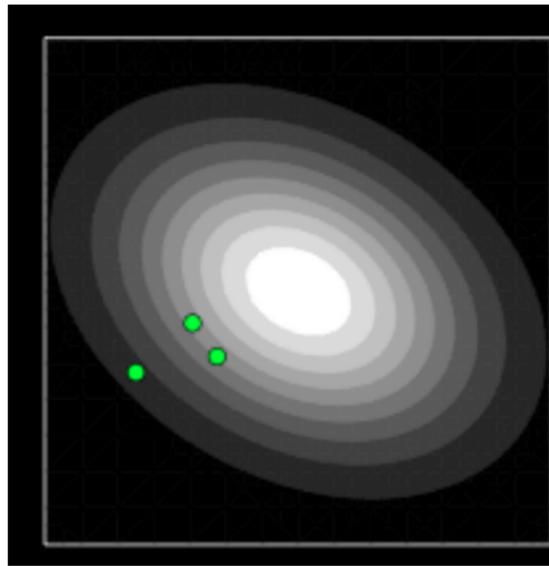
Courtesy of Sebastian Seehars



multiplicity + 1

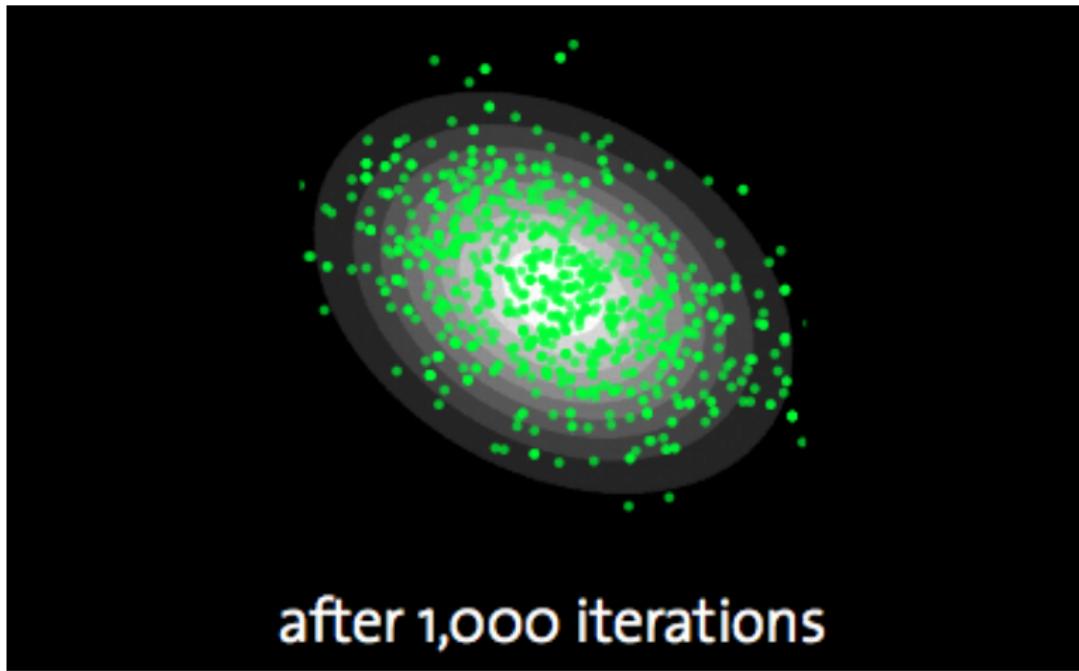
Metropolis-Hastings

Courtesy of Sebastian Seehars



Metropolis-Hastings

Courtesy of Sebastian Seehars



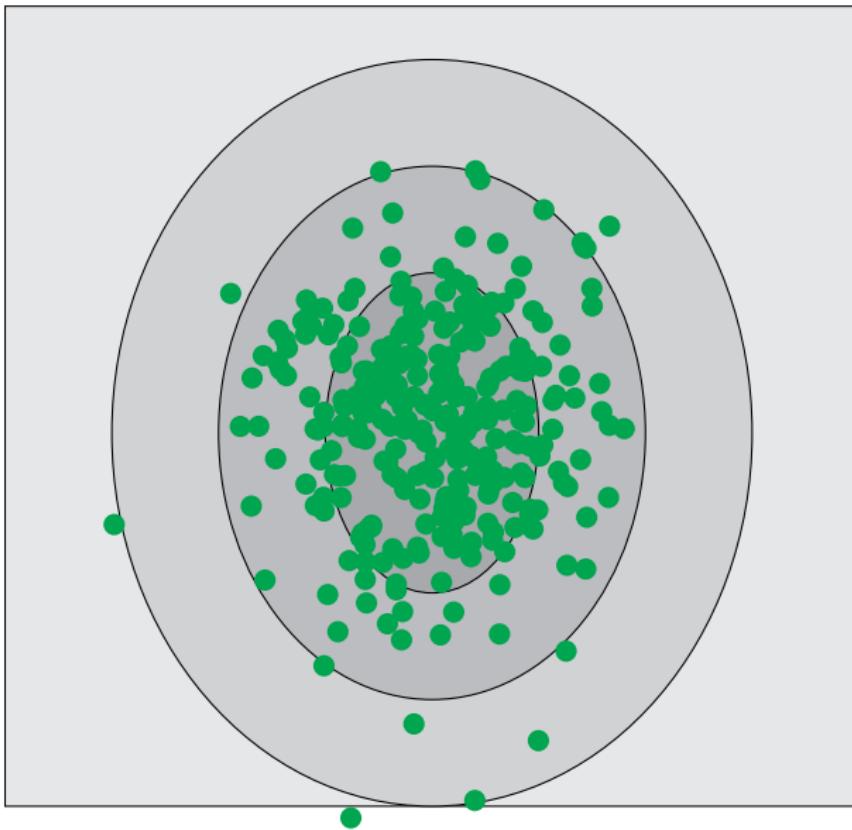
after 1,000 iterations

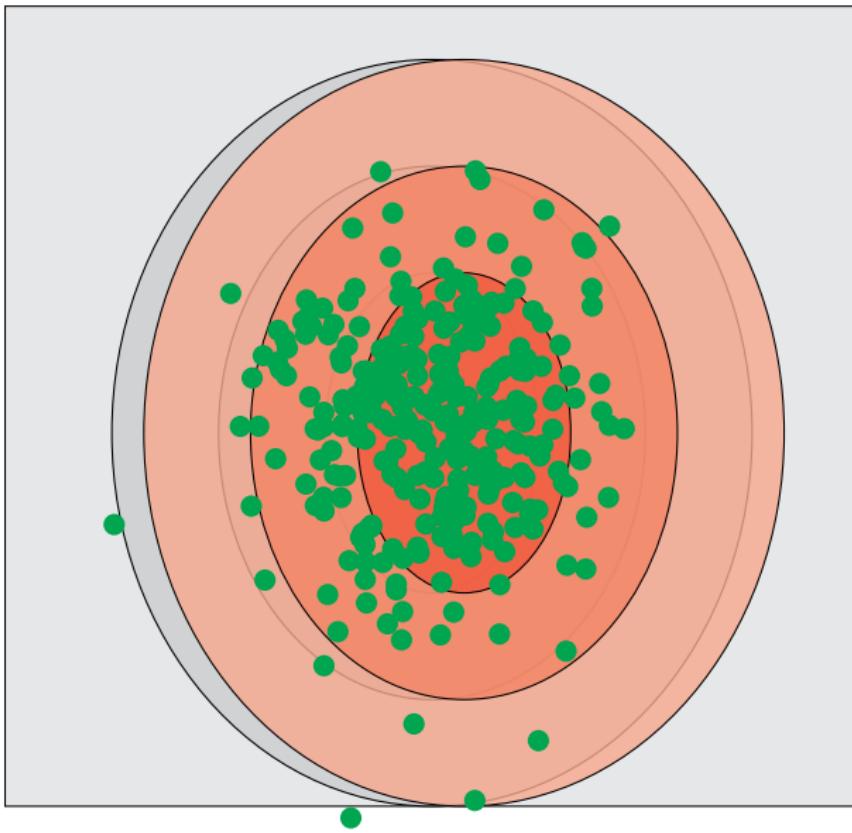
Importance Sampling

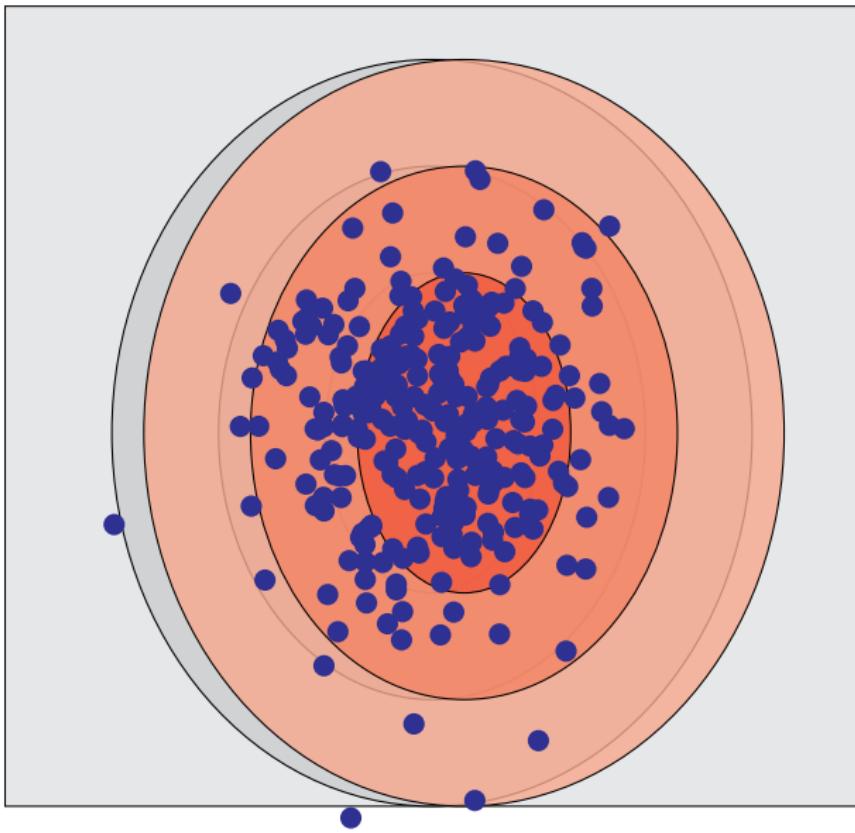
Not a real sampling

- start from an existing, **similar** distribution
- **add the new likelihood** at every already sampled point
- **write a new chain**, with the multiplicity changed:

$$\tilde{N} = N \frac{\tilde{\mathcal{L}}}{\mathcal{L}}$$







Benefits of Importance Sampling

When to use it?

- Existing run with a set of slow experiments
- Want to add one more experiment (prior on H_0)
- This should be very fast!